Codifix

Knowledge and governance for the Agentic Enterprise

Help teams find trusted expertise across boundaries and reuse what they learn — with local control.

Starting with incident handover between application and platform teams.

The incident moves. The context does not.

Illustrative incident · payment latency after a deployment

Application team

Knows the change and symptoms. Hands over a ticket and partial logs.

handover →

Platform team

Knows the infrastructure history. Rebuilds the missing context.

Repeated investigation

Interrupted experts

Corrections stay local

Pain and outcomes to validate in discovery; this is not a customer case study.

From agentic work to an Agentic Enterprise

01

Agentic Engineering

Agents help people build software.

02

Agentic OS

An environment for agentic work. Omarchy is an analogy.

03

Agentic Enterprise

Business units exchange selected expertise under local authority.

Broader vision: a network of business units. Current scope: governed knowledge exchange.

Every unit owns its knowledge and authority

Application business unit

Knowledge collections
Local agents
Independent administration

Authorized connection



Selected expertise
+ cited answers

Platform business unit

Knowledge collections
Local agents
Independent administration

ATProto provides identity and signed records; Memory governs access and private Q&A.

One handover becomes reusable expertise

1 Handover

Responder identifies the incident context.

2 Discover

Find permitted expertise across the connection.

3 Cited answer

Source team returns an answer with source citations.

4 Human action

Responder checks, acts, and records corrections.

5 Review at origin

The owning team approves the corrected knowledge.

6 Reuse later

A later eligible incident retrieves the correction.

Components implemented; end-to-end customer improvement still requires a measured pilot.

Measure better handovers, not activity volume

Faster correct action	Handover → first correct action and accepted resolution.
Fewer expert interruptions	Interruptions and expert minutes per eligible incident.
Learning that survives	Approved-correction reuse and recurrence of corrected errors.

Recovered capacity becomes cash savings only when spending actually changes.

Agree thresholds and comparable cases before declaring success. Retain failures and refusals.

Implemented software; customer proof still ahead

Capability	Evidence today	Remaining proof
Governed federation	Signed, reviewed knowledge exchange	Customer deployment and recovery gates
Private gateway Q&A	Bounded request/answer and source citations	Answer quality and operational impact
Origin-owned corrections	Versioned publication and re-review controls	Measured correction reuse in later incidents

Controlled demonstrations used synthetic data and isolated services on one host.

Implementation is not validated quality, customer readiness or a production deployment.

Land with two units. Expand through trusted links.

Economic buyer

CIO / CTO or technology-operations leadership

Operational champions

Application and platform leaders with named owners on both sides

Start

Two independently administered units

Expand internally

Divisions and acquired entities

Extend the network

Authorized external partners

Buyer and expansion logic are hypotheses to test; no customer or partner commitments claimed.

Compete on controlled exchange and ownership

Enterprise search & workplace agents

Glean and Rovo already offer connectors,
permission-aware access and answers.

Agent memory platforms

Mem0, Zep and Letta offer persistent memory;
shared memory and governance are established capabilities.

Codifix focus: exchange across independent administrations, with corrections reviewed and owned by the originating team.

Defensibility must be earned through workflow adoption, governed history and delivery experience.

An enterprise platform commitment

£80,000 / year

One enterprise account
Two independently administered deployments

Separate commercial components

Additional deployments
Onboarding
Usage overages

Replicated copies do not become separately billable new knowledge.

Pricing hypothesis. Initial Base includes no pilot, onboarding or overage receipts.

Founder-led discovery, then a qualified pilot

Qualify the buyer and boundary

Recurring pain · sponsor and budget access · independent control
Named owners · approved sharing · observable baseline

2 weeks

Baseline

Current handover workflow

4 weeks

Assisted operation

Comparable eligible incidents

2 weeks

Analysis

Quality, value and full cost

Commercial qualification cannot override deployment, security, recovery or evaluation gates.

Advance when the evidence supports it

1 **Deployment readiness**

Security, authorization, recovery and evaluation gates.

2 **Measured pilot outcome**

Correctness, handover time, interruptions and correction reuse.

3 **Validated delivery economics**

Observed inference, support, curation and assurance costs.

4 **Repeatable onboarding**

Independent teams onboard within measured effort and limits.

Readiness shifts the pipeline. There is no unconditional GA date or committed customer count.

Founder-led, with explicit assurance ownership

Olivier Van Acker

Founder · creator of Memory Layer

Built the Rust memory platform, code-graph retrieval, MCP integration and evaluation tooling.

Planned organization

Nine employees, including founder,
+ fractional counsel
Engineering, deployment, commercial, support and an assurance lead.

Assurance lead owns the programme:
ISO 27001 · ISO 42001 · SOC 2 Type II

Employee AI per person, per month:
£5k allowance · £2.5k expected spend.

Hiring, AI utilization and throughput are hypotheses. Certifications are not completed.

Fund the next proof points

Capital should demonstrate

A trustworthy deployment	Pass security, recovery and evaluation gates.
A valuable workflow	Measure quality and operational outcomes in a qualified pilot.
A repeatable business	Validate delivery cost, packaging and onboarding effort.

Discuss milestones and financing structure.

Round size follows the evidence, scope and runway required. No fixed fundraising ask.

Provisional financial scenarios in the appendix

codifix

Appendix · architecture and trust boundaries

AT Protocol foundation

Stable identity (DIDs)
Lexicon record schemas
Signed repositories and sync

Memory governance

Authorization and publication history
Reviewed, versioned knowledge
Private Q&A through the gateway

Trust and disclosure limits

Gateway availability and enforcement are trusted dependencies.
Authorized recipients can read and retain copies; revocation cannot recall bytes.
Signed records prove authenticity, not the truth of the content.

Native Spaces and broader delegated execution remain future work.

Appendix · evidence without a quality claim

60

Knowledge answers
awaiting genuine blinded review

0 / 15

Successful complete
coding sequences

R3 development run · 17 September 2026

All 75 executions completed. Harness validity and quality did not pass.
No validated quality improvement, perfect recall or token-saving claim.

Controlled federation demonstrations establish bounded behavior, not customer ROI.

Appendix · Base commercial forecast

Shared economics · GBP	Year 1	Year 2	Year 3
Year-end ARR	£160k	£900k	£2.24m
Recognized revenue	£42k	£552k	£1.60m
Paying accounts (expected)	2.00	10.40	22.38
London / NY gross margin	−80.8% / −101.4%	52.1% / 48.0%	74.2% / 72.7%

Pilot funnel: starts from month 5, two pilot months, 50% conversion.
Net revenue retention: (1 − 5% churn) × (1 + 45% expansion) = 137.75%.

36 months from 1 October 2026. Fractional accounts are expected values, not signed customers.

Appendix · alternative operating bases

24-month funding objective	London · GBP	New York · USD
Gross funding requirement	£3.23m	$4.91m
Expected employee AI spend	£420k	$566k
AI headroom retained as cash	£420k	$566k
M25–27 operating buffer	£504k	$744k

New York gross funding in common GBP: £3.65m at $1.3476/£.

Fund to month 24; retain AI headroom plus months 25–27 operating outflows.
Alternative bases, never combined. External assurance costs remain unquoted.

AI headroom is cash, not a second expense. These scenarios do not determine a round size.